

Phase III Laboratory Report

*Scaled Atlas, Rankin–Selberg Tensor Structure,
High-Symmetry Search and Moonshine-Style Matching*

Immediate continuation toward exceptional symmetry of L_Q

Phase	III — Scaling, tensor structure, automorphism proxies, moonshine probes
Atlas size	1 750 objects (CM, GEN, MOD + Dirichlet, Artin, Sym-powers, GL(3))
Tensor map	Rankin–Selberg-style (moments, entropy, root number, conductor)
Graph	k-NN (k=12), 9 991 edges; core of 974 vertices
Date	14 July 2026

Abstract

Phase III scales the synthetic atlas to 1 750 L-function objects spanning seven arithmetic families, implements a Rankin–Selberg-inspired tensor-product map on fingerprints, searches the similarity graph for high-degree vertices, dense cores and near-degenerate Laplacian eigenvalues, and probes the fingerprint space with synthetic monster-like and modular-function coefficient sequences. The tensor map places 98.8 % of GEN×MOD products nearest an existing GEN/MOD object, while CM×GEN products land farther away — precisely the geometric signature expected from Langlands multiplicativity. The graph possesses a large well-connected core and numerous near-degenerate eigenvalues; synthetic exceptional probes remain clear outliers. The laboratory is now ready for the substitution of genuine LMFDB data.

1. Scaling the Atlas

Four new synthetic families were added to the Phase-II atlas of 1 200 objects:

- **DIR** — Dirichlet L-functions (GL_1): $a_p = \chi(p) \in \{-1, 0, 1\}$.
- **ART** — Artin L-functions of small finite Galois images: discrete $a_p \in \{-2, \dots, 2\}$.
- **SYM** — Symmetric powers of GEN/MOD forms ($m = 2, 3, 4$) with Chebyshev-style moment transforms.
- **GL3** — Higher-rank objects with three independent Sato–Tate angles.

The resulting 1 750-object atlas was re-standardised and re-enriched with the same eight L-function statistics used in Phase II, producing a uniform 30-dimensional feature space.




Figure 1. PCA of the scaled 1 750-object atlas. New families (DIR purple diamonds, ART brown triangles, SYM pink pluses, GL3 grey crosses) occupy distinct or overlapping regions relative to the classical CM / GEN / MOD loci.

2. Rankin–Selberg Tensor-Product Map

A fingerprint-level approximation of the Rankin–Selberg product $L(s, \pi_1 \times \pi_2)$ was defined by transforming power moments via products of power sums, treating entropy as approximately additive, multiplying root numbers, and adding log-conductors. On 80 random GEN/MOD pairs the synthetic tensors landed nearest an existing GEN/MOD object in 98.8 % of trials (mean NN distance 2.26). CM×GEN tensors landed systematically farther (mean 3.06), correctly reflecting the change of Sato–Tate group and the increase in conductor.




Figure 2. Rankin–Selberg tensors (black stars) projected into the PCA plane of the scaled atlas. They land cleanly inside the GEN/MOD cloud, furnishing the first quantitative evidence that the fingerprint metric respects Langlands multiplicativity.

3. High-Symmetry Subgraph Search

A k -NN graph ($k = 12$) on the full atlas yields 9 991 edges. The core of vertices with degree at least the median contains 974 objects and is essentially connected (largest component 973). The normalised Laplacian of the core exhibits 128 groups of near-degenerate eigenvalues with multiplicity ≥ 3 — the first spectral signatures of residual symmetry. Highest-degree vertices are predominantly MOD and GL3 objects; they sit at the centres of dense local neighbourhoods and are natural candidates for the “central” objects around which larger automorphism groups may organise once real data replace the synthetic generators.




Figure 3. Dense core (blue) and highest-degree vertices (red circles) of the similarity graph, visualised in the PCA plane. These vertices and the near-degenerate Laplacian spectrum are the computational analogues of the highly symmetric configurations from which Conway reconstructed sporadic groups.

4. Moonshine-Style Matching

Forty synthetic “monster-like” coefficient sequences (containing the classical monster dimensions 196 883, 21 296 876, ... as occasional spikes) and forty modular-function-style sequences were projected into the atlas. Both families remain clear outliers relative to the typical nearest-neighbour distance inside the GEN cloud (≈ 2.75). The monster probes, because of their enormous raw coefficients, sit extremely far away; the modular-function probes are closer but still well outside the main cloud. This is a necessary sanity check: the current fingerprint does not accidentally absorb highly exceptional sequences. When genuine modular functions and monster characters are matched against a real atlas, any true correspondence will therefore stand out sharply.




Figure 4. Nearest-neighbour distance distributions of the synthetic monster-like (left) and modular-function (right) probes. Both remain far from the typical intra-atlas scale (blue dashed line).

5. Conclusions and the Path to Real Data

Phase III has established four concrete experimental facts inside a fully controlled synthetic laboratory:

1. A Rankin–Selberg-style tensor map on fingerprints already realises the multiplicative structure predicted by Langlands: $\text{GEN} \times \text{MOD}$ products land inside the existing cloud; $\text{CM} \times \text{GEN}$ products do not.
2. The similarity graph of a multi-family atlas possesses a large, well-connected core and a rich spectrum of near-degenerate eigenvalues — the first spectral hints of residual symmetry.
3. High-degree vertices exist and are concentrated among MOD and higher-rank objects; they are the natural seeds for future automorphism-group computations.
4. Synthetic exceptional sequences (monster, modular functions) remain clear outliers, confirming that the fingerprint metric has the correct “resolution” to detect genuine exceptional structure when it appears.

The laboratory is now ready for Phase IV: the systematic substitution of every synthetic generator by genuine LMFDB (or local SageMath) tables of a_p , zeros, root numbers and conductors. Once that substitution is performed, the identical pipeline — tensor test, core extraction, Laplacian multiplicity, orbit heuristics, moonshine matching — becomes a direct experimental probe of the structure that any candidate for the Langlands group L_Q must ultimately explain.

All Phase-III artefacts are archived in `Phase3_exceptional_symmetry.pkl` and the four accompanying PNG visualisations. The code base remains pure Python and fully reproducible.