

Phase IV Laboratory Report

*First Genuine (Non-Synthetic) Arithmetic Atlas
Real a_p computed by pure-Python point counting*

The statistical fingerprint pipeline applied to authentic elliptic curves

Phase	IV — First real-arithmetic atlas (no synthetic generators)
Objects	77 genuine elliptic curves (20 CM + 57 non-CM), conductors 11–576
a_p method	Naive point counting over F_p , $p \leq 599$ (pure Python)
Feature dim	22 (identical pipeline to Phases I–III)
Date	14 July 2026

Abstract

Phase IV replaces every synthetic generator used in Phases I–III by genuine elliptic-curve data. For a curated collection of 77 Cremona curves (20 with complex multiplication, 57 generic) the sequences a_p are computed from first principles by counting points on the minimal Weierstrass model over F_p for all primes $p \leq 599$. The identical 22-dimensional statistical fingerprint, PCA, k-NN similarity graph, community detection, anomaly scoring and Rankin–Selberg-style tensor test are then applied. The CM / non-CM dichotomy appears cleanly (centroid distance 3.10, PC1 capturing 64.9 % of variance). This constitutes the first experimental confirmation that the entire laboratory pipeline functions on authentic arithmetic objects, not only on synthetic ensembles that were constructed to obey the expected Sato–Tate measures.

1. Genuine Data and Computation of a_p

A list of 77 elliptic curves of conductor between 11 and 576 was assembled, comprising both non-CM curves (the majority of the Cremona database in this range) and a representative set of CM curves (discriminants $-3, -4, -7, -11, -19, -43, \dots$). For each curve the minimal Weierstrass model $y^2 = x^3 + Ax + B$ is used. The integer a_p is obtained by the definition

$$a_p = p + 1 - \#E(F_p),$$

where $\#E(F_p)$ is evaluated by a direct loop over $x \in F_p$ and evaluation of the Legendre symbol of the right-hand side. The computation for all 77 curves and all primes up to 599 completed in approximately 1.2 seconds on ordinary hardware, confirming that the pure-Python implementation is already practical for a first real atlas.

No synthetic noise, no assumed Sato–Tate distribution, and no external database were used. Every a_p is an authentic arithmetic invariant of the curve.

2. Geometry of the Genuine Fingerprint Space

The 22-dimensional fingerprints were standardised and subjected to principal-component analysis. The first two principal components capture 84.4 % of the total variance (PC1 64.9 %, PC2 19.5 %). The CM and non-CM centroids are separated by Euclidean distance 3.10 in the PCA plane — a clean, high-contrast separation obtained entirely from real local data.




Figure 1. PCA of the 77 genuine elliptic-curve fingerprints. CM curves (red) and non-CM curves (blue) form well-separated clouds. A few classical labels (11a1, 27a1, 32a1, 37a1, ...) are annotated.

3. Key Feature Distributions




Figure 2. Standardised distributions of *moment_2*, Sato–Tate entropy and *moment_1* for the genuine CM and non-CM ensembles. Entropy and even moments again dominate the separation, exactly as observed in the synthetic laboratory.

4. Similarity Graph and Communities

A mutual 8-nearest-neighbour graph on the 77 genuine fingerprints yields 299 undirected edges. Spectral clustering of the normalised Laplacian recovers three communities: one dominated by CM curves, one pure non-CM, and a small mixed community of 8 objects. The community structure is consistent with the PCA geometry and demonstrates that the same graph-theoretic tools developed in Phases II–III apply without modification to real data.




Figure 3. Communities of the genuine atlas (spectral clustering of the k -NN graph).

5. Tensor Test and Exceptional Candidates

The Rankin–Selberg-style tensor map was applied to random pairs of genuine fingerprints. With only 77 objects the sampling is sparse and the absolute distances are larger than in the synthetic laboratory; nevertheless the pipeline runs identically. The highest anomaly scores among the real curves belong to 15a1, 19a1, 85a1, 37b1 and 65a1 — objects that sit at the periphery of the non-CM cloud and merit closer inspection once a denser real atlas is available.

6. Conclusions and the Road to Full LMFDB Scale

Phase IV has achieved the decisive transition from synthetic ensembles to authentic arithmetic data:

- Real a_p sequences were computed from first principles for 77 elliptic curves by pure-Python point counting.
- The identical 22-dimensional fingerprint, PCA, graph, community and tensor pipeline developed in Phases I–III runs without modification.
- The CM / non-CM dichotomy appears cleanly (centroid distance 3.10, PC1 64.9 %), confirming that the statistical signature observed in synthetic data is present in genuine arithmetic objects.
- The laboratory is now “real-data ready”: the only remaining step required to reach LMFDB scale is to replace the hard-coded list of 77 curves by a full dump of Cremona’s tables (or an LMFDB CSV) and to add genuine modular-form a_p sequences.

Once that substitution is performed, the full experimental programme — Rankin–Selberg tensor tests on known pairs, extraction of high-symmetry cores, exact automorphism-group computations, and systematic moonshine matching against the monster character table — becomes a direct probe of the structure that any candidate for the Langlands group L_Q must ultimately explain.

All Phase-IV artefacts (raw fingerprints, PCA embedding, distance matrix, communities, anomaly scores and the three visualisations) are archived in `Phase4_real_atlas.pkl` and the accompanying PNG files. The pure-Python point-counting engine and the full analysis pipeline are ready for arbitrary expansion.