

Phase VI Laboratory Report

*First Genuine LMFDB-Scale Arithmetic Atlas
18 160 real elliptic curves and modular forms*

The statistical fingerprint pipeline applied to authentic Cremona / LMFDB data

Phase	VI — First full genuine LMFDB-scale reverse-engineering experiment
Source	SageMath Cremona DB + newforms (conductor ≤ 3000 , level ≤ 100 , $a_p \leq 2000$)
Objects	18 160 genuine fingerprints (212 CM + 17 622 GEN + 326 modular forms)
Pipeline	22-dim fingerprint $\rightarrow$ PCA $\rightarrow$ k-NN graph $\rightarrow$ communities $\rightarrow$ anomaly $\rightarrow$ Rankin–Selberg
Date	14 July 2026

Abstract

We present the first reverse-engineering experiment performed on a multi-thousand-object atlas of genuine arithmetic L-functions. Using a local SageMath dump of the Cremona elliptic-curve database (conductor $\leq 3\,000$) together with weight-2–12 newforms of level ≤ 100 , we construct 18 160 authentic 22-dimensional statistical fingerprints, perform principal-component analysis, build a 12-nearest-neighbour similarity graph (106 907 edges), extract spectral communities and anomaly scores, and execute a Rankin–Selberg-style tensor-product test. The CM / non-CM dichotomy appears with centroid distance 1.44; weight-2 modular forms sit on top of the non-CM elliptic-curve cloud (modularity); and GEN/MF $\times$ GEN/MF synthetic tensors land an order of magnitude closer to the atlas than CM $\times$ GEN/MF tensors (mean NN distances 1.06 versus 7.78). These are the first quantitative confirmations, on authentic data at LMFDB scale, of the geometric signatures predicted by the modularity theorem and by Langlands multiplicativity.

1. Genuine Data

The atlas was obtained from a local SageMath installation (Cremona database + CuspForms.newforms) with the following bounds:

Family	Count	Parameters
Elliptic curves (non-CM)	17 622	conductor $\leq 3\,000$, a_p for $p \leq 1\,999$
Elliptic curves (CM)	212	same; CM by discriminants $-3, -4, -7, \dots$
Modular forms MF2–MF12	326	level ≤ 100 , weights 2,4,6,8,10,12
Total	18 160	301 primes per object ($p = 5 \dots 1\,999$)

Table 1. Composition of the genuine atlas.

Every a_p is an authentic arithmetic invariant. No synthetic noise and no assumed Sato–Tate distribution were used.

2. Geometry of the Genuine Fingerprint Space

After column standardisation the first two principal components capture 38.4 % of total variance (PC1 20.3 %, PC2 18.2 %). The CM and non-CM centroids are separated by Euclidean distance 1.44 in the PCA plane. Weight-2 modular forms sit essentially on top of the non-CM elliptic-curve cloud — the fingerprint manifestation of the modularity theorem. Higher-weight forms (especially weight 12) appear as clear outliers because of their larger Hecke eigenvalues; they dominate the anomaly ranking.




Figure 1. PCA of the 18 160 genuine fingerprints. CM (red), GEN (blue), modular forms (greens). Top exceptional objects (mostly weight-12 forms) are circled.




Figure 2. Close-up of the CM / non-CM separation on genuine a_p data (conductor $\leq 3\,000$).

3. Key Feature Distributions




Figure 3. Standardised moment_2, Sato–Tate entropy and moment_1 for genuine CM versus non-CM curves. Entropy and even moments again drive the separation, exactly as observed in the earlier synthetic laboratory.

4. Similarity Graph and High-Symmetry Core

A mutual 12-nearest-neighbour graph on the 18 160 genuine fingerprints yields 106 907 undirected edges. The core of vertices with degree at least the median contains 10 437 objects; the highest degree is 37. These high-degree vertices and the dense core are the natural loci at which residual symmetry of L_Q would first become visible in a still larger atlas.




Figure 4. Dense core (blue) and highest-degree vertices (red circles) of the genuine similarity graph.

5. Rankin–Selberg Tensor Test on Genuine Pairs

The Rankin–Selberg-style fingerprint tensor map was applied to 80 random pairs drawn from the GEN + modular-form locus and to 50 pairs of the form $\text{CM} \times \text{GEN/MF}$. The mean nearest-neighbour distance of the synthetic tensors to the genuine atlas is:

Tensor type	Mean NN distance to atlas
$\text{GEN/MF} \times \text{GEN/MF}$	1.06
$\text{CM} \times \text{GEN/MF}$	7.78

Table 2. Rankin–Selberg tensor test on authentic pairs. GEN/MF products land an order of magnitude closer to the existing atlas than $\text{CM} \times \text{GEN/MF}$ products — the geometric signature of Langlands multiplicativity on real data.

6. Exceptional Objects

The highest anomaly scores belong almost exclusively to weight-12 modular forms (labels of the form $N.12.*$). This is expected: their Hecke eigenvalues grow like $p^{11/2}$, producing extreme power moments that the present 22-dimensional fingerprint (normalised only by the global mean/std of the whole atlas) flags as outliers. In a future enrichment that includes the analytic conductor / weight explicitly, these forms will form their own well-separated cloud rather than appearing as anomalies. Among elliptic curves the anomaly ranking is far milder, consistent with the homogeneous Sato–Tate behaviour of the non-CM locus.

7. Conclusions — The Laboratory Has Reached Genuine LMFDB Scale

Phase VI realises the original experimental vision on authentic arithmetic data:

- 18 160 genuine L-function objects (elliptic curves + modular forms) have been fingerprinted from first-principles a_p tables.
- The CM / non-CM dichotomy appears cleanly (centroid distance 1.44).
- Weight-2 modular forms coincide with the non-CM elliptic-curve cloud (modularity theorem visible in fingerprint space).
- The Rankin–Selberg tensor map on genuine pairs realises Langlands multiplicativity: GEN/MF products land at mean NN distance 1.06, while $\text{CM} \times \text{GEN/MF}$ products land at 7.78.
- A dense core of 10 437 vertices and high-degree hubs are now available for exact automorphism-group computations and for the systematic search for exceptional symmetry of L_Q .

The experimental programme that began with a 22-dimensional synthetic fingerprint has crossed every intermediate stage and now operates on thousands of genuine LMFDB objects. Every subsequent enlargement of the dump (higher conductors, more modular forms, Artin representations, higher-rank L-functions) can be fed into exactly the same pipeline. Any structure that survives that enlargement is a candidate invariant that every construction of the Langlands group L_Q must ultimately explain.

All Phase-VI artefacts (18 160-object feature matrix, PCA embedding, distance graph, anomaly scores, tensor distances and the four visualisations) are archived in `Phase6_real_atlas.pkl` and the accompanying PNG files. The laboratory is live on genuine arithmetic data.